

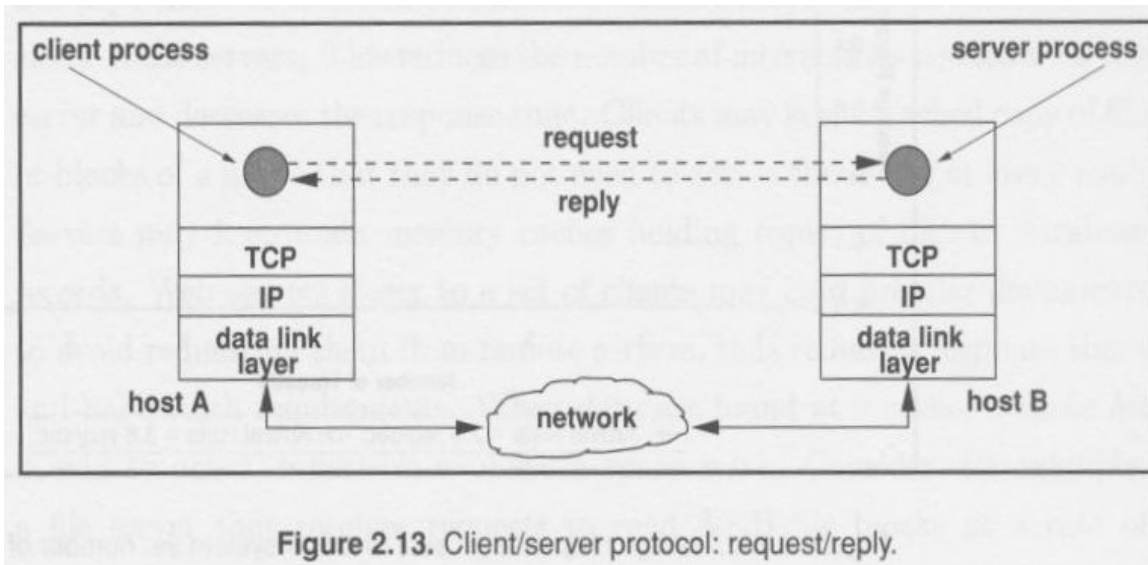
Network Programming

Client Server Paradigm


- ***Server application is “listener”***
- Waits for incoming message
- Performs service
- Returns results
- ***Client application establishes connection***
- Sends message to server
- Waits for return message

Client-Server Communication

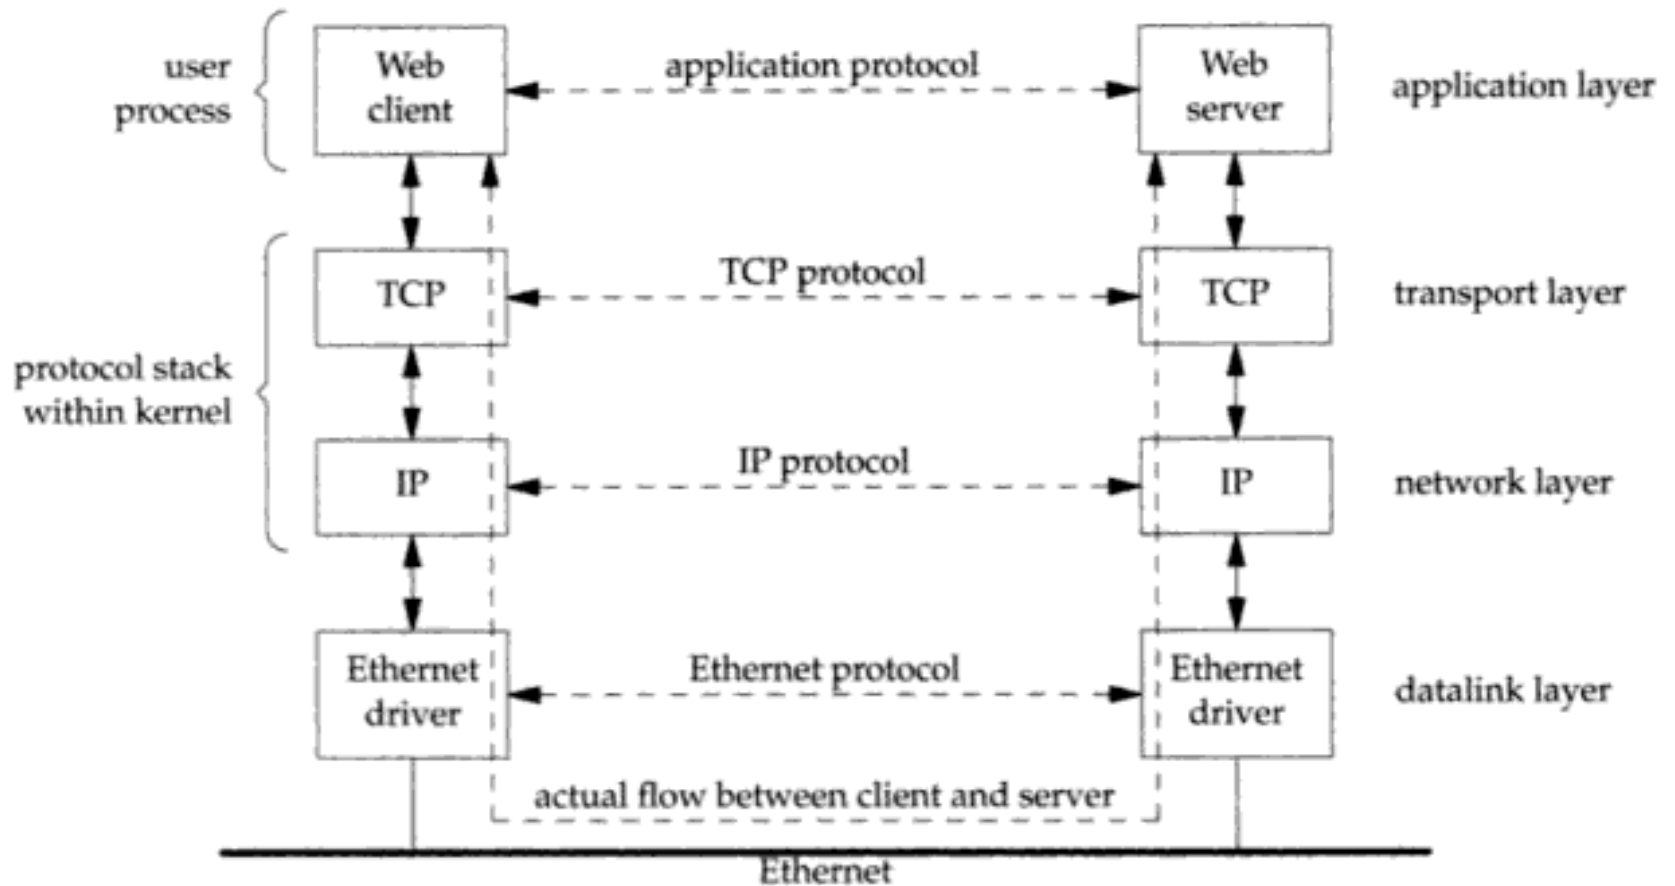
- Clients and servers exchange messages through transport protocols; e.g., TCP or UDP
- Both client and server must have same protocol stack and both interact with transport



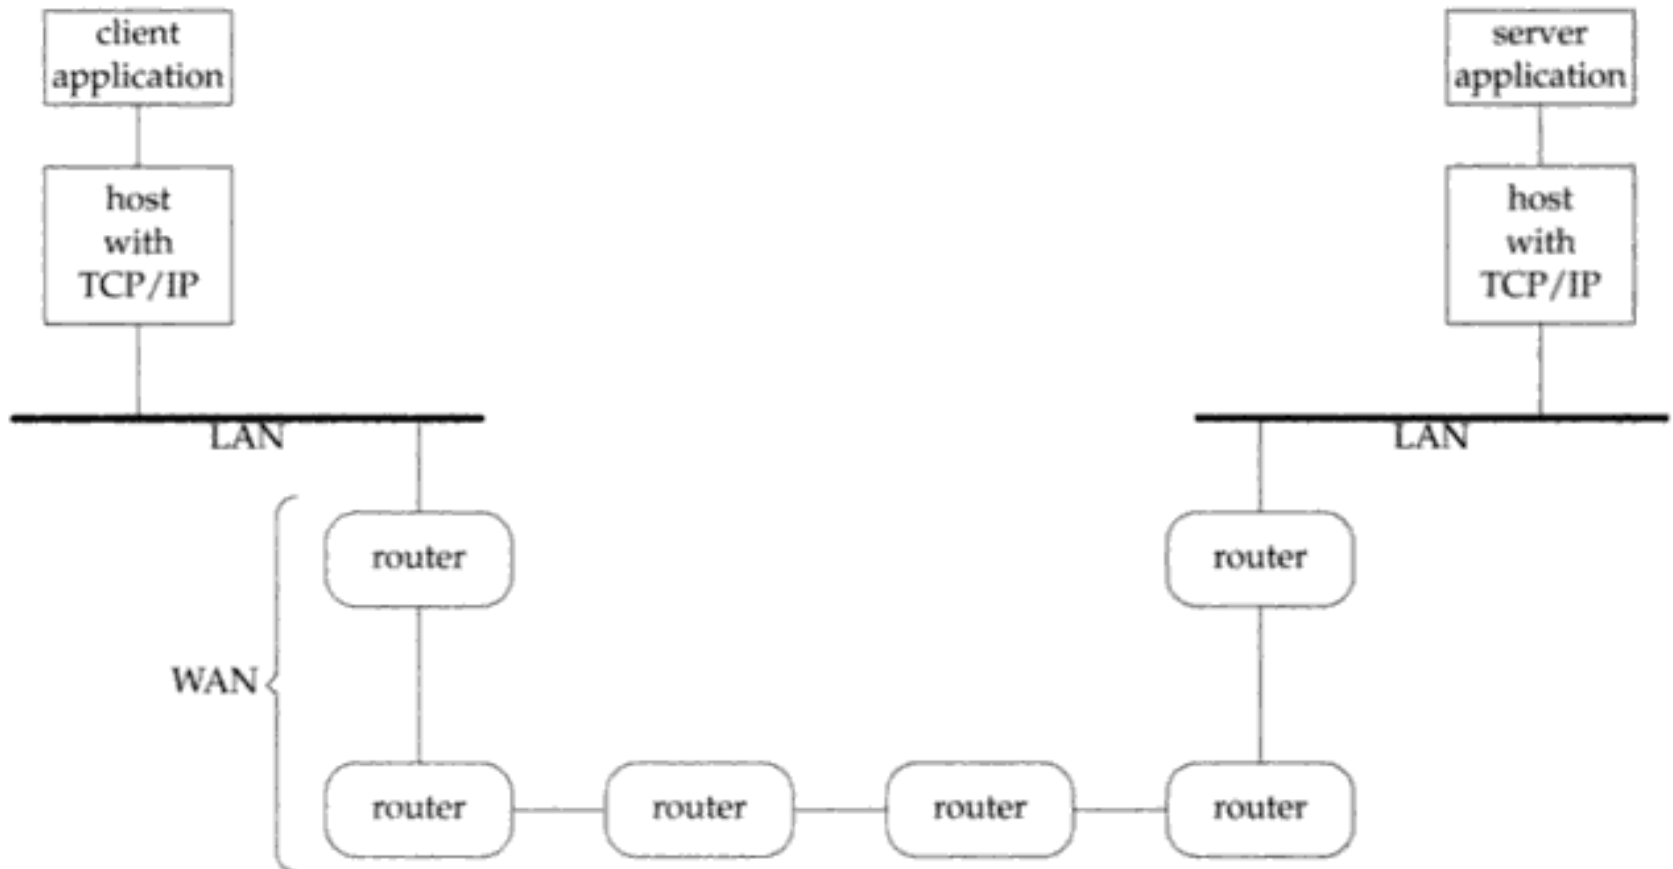
Interaction Between Client-Server



C S on same Ethernet



C S on different LAN



OSI Layers In Client Server

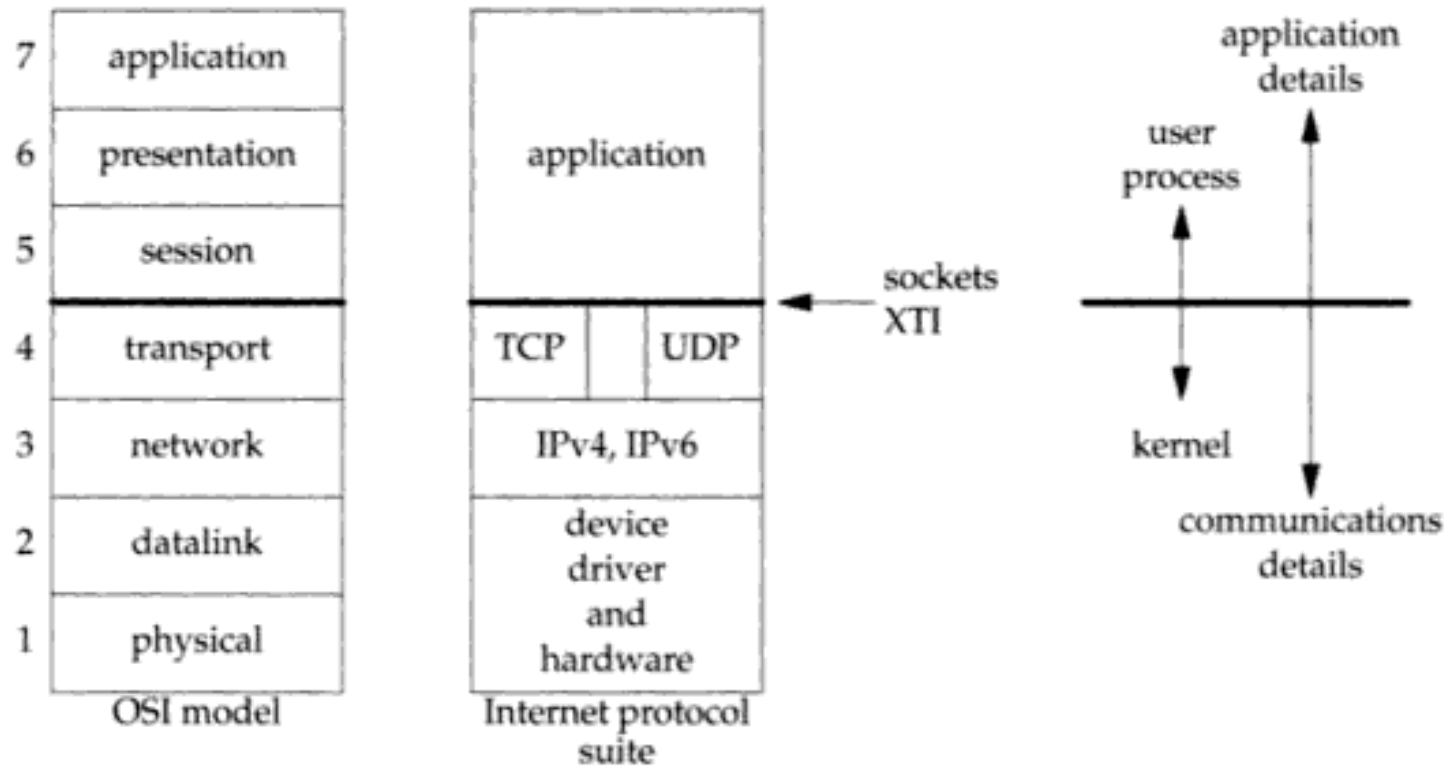


Figure 1.14 Layers in OSI model and Internet protocol suite.

Programming Interfaces: Sockets

Client-Server Communication Model

□ Service Model

❖ Concurrent:

- Server processes multiple clients' requests simultaneously

❖ Sequential:

- Server processes only one client's requests at a time

❖ Hybrid:

- Server maintains multiple connections, but processes responses sequentially

□ Client and server categories are not disjoint

- ❖ A server can be a client of another server
- ❖ A server can be a client of its own client

Sockets

```
int sockfd = socket(PF_INET, SOCK_STREAM, 0);
```

- ❑ Creates a new socket
- ❑ SOCK_STREAM = TCP
- ❑ SOCK_DGRAM = UDP

Socket Address in General

Socket Addresses

- Need to be able to specify addresses
- Clients needs to
 - Specify the IP address of the destination
 - Specify the port number of the destination
 - (Obtain a local port number)
- Servers need to
 - Associate the (well-known) server port number to the socket
 - (Restrict the service to particular IP addresses)
- Addresses are specified in network byte order (which is big endian)

Sockaddr struct

IP Socket Addresses

```
struct sockaddr_in {
    sa_family_t    sin_family;    /* address family: AF_INET */
    u_int16_t      sin_port;      /* port in network byte order */
    struct in_addr  sin_addr;      /* internet address */
};

/* Internet address. */
struct in_addr {
    u_int32_t      s_addr;        /* address in network byte order */
};
```

- Small variations between different operating systems
- `inet_addr()` can be used to convert a dotted-quad string to a binary representation
 - in network byte order

```
in_addr_t inet_addr(const char *cp);
```

Socket Address

Requesting Socket Addresses

```
int  getsockopt(int  s, int level, int optname, void *optval,  
socklen_t *optlen);  
  
int  setsockopt(int s, int level,  int  optname,  const void  
*optval, socklen_t optlen);
```

- *Getpeername*
 - Get name (socket address) of the peer connected to the socket
 - Remote side
- *Getsockname*
 - Get name (socket address) of the socket
 - . . .

SOCKET

Sockets Are for Networks in General...

```
int socket(int domain, int type, int protocol);
```

<i>Domain (family)</i>	<i>Purpose</i>	<i>Man page</i>
PF_UNIX, PF_LOCAL	Local communication	unix(7)
PF_INET	IPv4 Internet protocols	ip(7)
PF_INET6	IPv6 Internet protocols	
PF_IPX	IPX - Novell protocols	
PF_NETLINK	Kernel user interface device	netlink(7)
PF_X25	ITU-T X.25 / ISO-8208 protocol	x25(7)
PF_AX25	Amateur radio AX.25 protocol	
PF_ATMPVC	Access to raw ATM PVCs	
PF_APPLETALK	Appletalk	ddp(7)
PF_PACKET	Low level packet interface	packet(7)
<i>Type</i>	<i>Purpose</i>	
SOCK_STREAM	Sequenced, reliable, two-way, connection-based byte streams.	
SOCK_DGRAM	Datagrams.	
SOCK_SEQPACKET	Sequenced, reliable, two-way connection-based data transmission path for datagrams	
SOCK_RAW	Raw network protocol access.	
SOCK_RDM	A reliable datagram layer that does not guarantee ordering.	
SOCK_PACKET	Obsolete and should not be used in new programs	
<i>Protocol</i>	<i>Purpose</i>	
IPPROTO_TCP	TCP (Only legal protocol for PF_INET + SOCK_STREAM)	
IPPROTO_UDP	UDP (Only legal protocol for PF_INET + SOCK_DGRAM)	
0	Default protocol for given domain and type	

Socket Options

Socket Options

<i>SOL_SOCKET</i>	<i>Purpose</i>
SO_SNDBUF	Maximum send buffer size
SO_RCVBUF	Maximum receive buffer size
SO_KEEPALIVE	Enable/disable keep-alive messages (connection-oriented sockets)
SO_BROADCAST	Allow sending/reception of broadcast packets (datagram sockets)
SO_REUSEADDR	Allow reuse of local address in bind calls
<i>SOL_IP</i>	<i>Purpose</i>
IP_TTL	TTL field for packets sent
IP_ADD_MEMBERSHIP	Join a multicast group
IP_DROP_MEMBERSHIP	Leave a multicast
IP_MTU	Get current path MTU (connected sockets)
<i>SOL_TCP</i>	<i>Purpose</i>
TCP_MAXSEG	Maximum segment size
TCP_NODELAY	Disable the Nagle algorithm

- See "man 7 socket", "man 7 ip", "man 7 tcp", ...

Byte Ordering

□ Big Endian vs. Little Endian

❖ Little Endian (Intel, DEC):

- Least significant byte of word is stored in the lowest memory address

❖ Big Endian (Sun, SGI, HP):

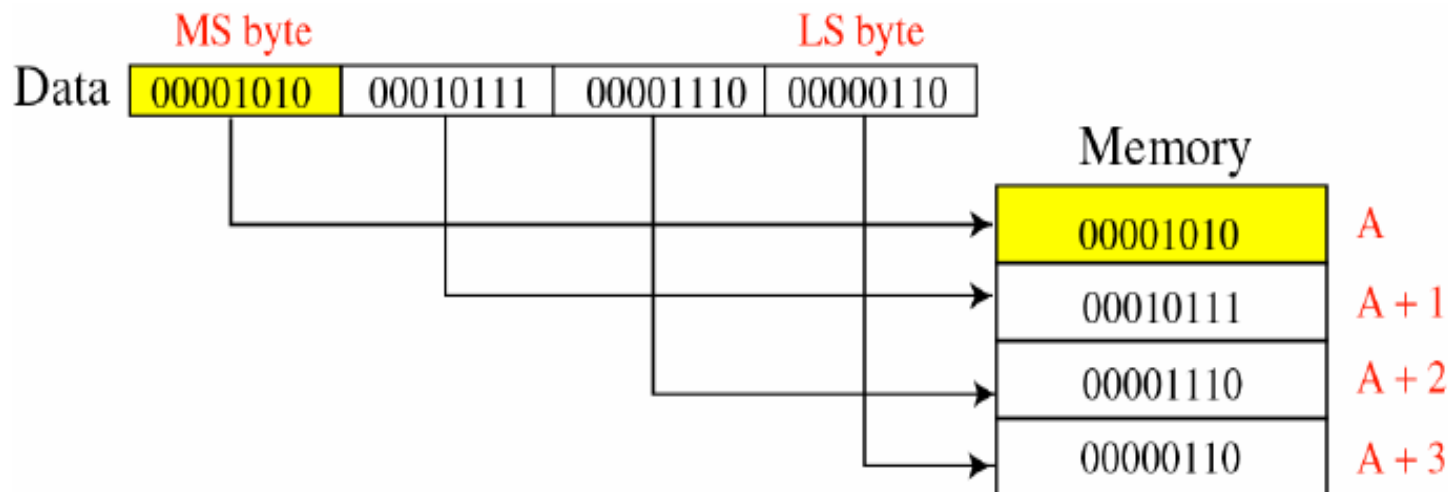
- Most significant byte of word is stored in the lowest memory address

❖ Network Byte Order = Big Endian

- Allows both sides to communicate
- Must be used for some data (i.e. IP Addresses)
- Good form for all binary data

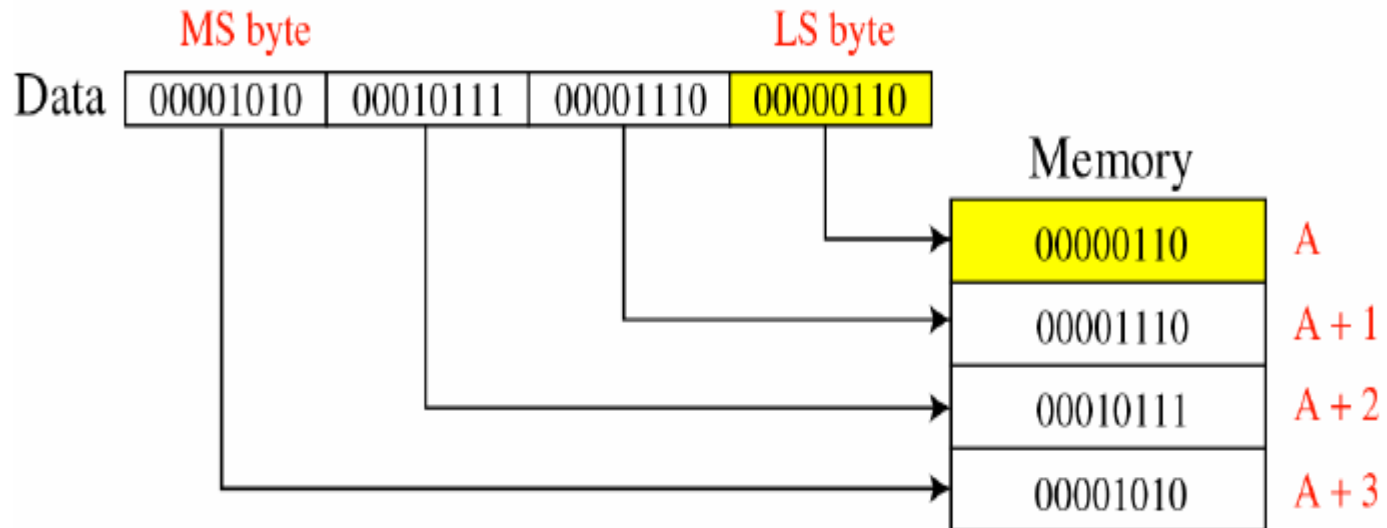
Big-Endian Format

Big-Endian Byte Order



Little Endian Format


Little-Endian Byte Order



Network Transformation

Byte-Order Transformation

- "short"—16 bits
- "long"—32 bits



Byte Ordering Functions

- ❑ 16- and 32-bit conversion functions (for platform independence)
- ❑ Examples:

```
int m, n;  
short int s, t;
```

<code>m = ntohl (n)</code>	net-to-host long (32-bit) translation
<code>s = ntohs (t)</code>	net-to-host short (16-bit) translation
<code>n = htonl (m)</code>	host-to-net long (32-bit) translation
<code>t = htons (s)</code>	host-to-net short (16-bit) translation

Connecting the socket

```
connect(sockfd, (struct sockaddr *) their_addr, sizeof their_addr);
```

- ❑ Actually creates a TCP connection, contacting a server
- ❑ Afterwards, use file descriptor like a regular file - read, write, close
- ❑ Return values (Error):
 - ❖ ETIMEDOUT: Client TCP does not receive a response to the SYN packet
 - ❖ ECONNREFUSED: No process is waiting on the server at the port specified
 - ❖ EHOSTUNREACH: If the client's SYN packet gets an ICMP "destination unreachable" message

bind

```
bind(sockfd, (struct sockaddr *) my_addr, sizeof my_addr);
```

- ❑ Connects a socket with a well-known *incoming* port on a local system
- ❑ Common error value:
 - ❖ EADDRINUSE – Address / port already in use

Listen

```
listen(sockfd, 10);
```

- ❑ Start accepting connections
- ❑ 10 is the amount of backlog connections
 - ❖ If more than 10 clients waiting, computer will refuse new connections
- ❑ How is the backlog arrived at?

Backlog entries????

- ❑ Backlog = Entries in incomplete connection queue + completed connection queue

Figure 4.7. The two queues maintained by TCP for a listening socket.

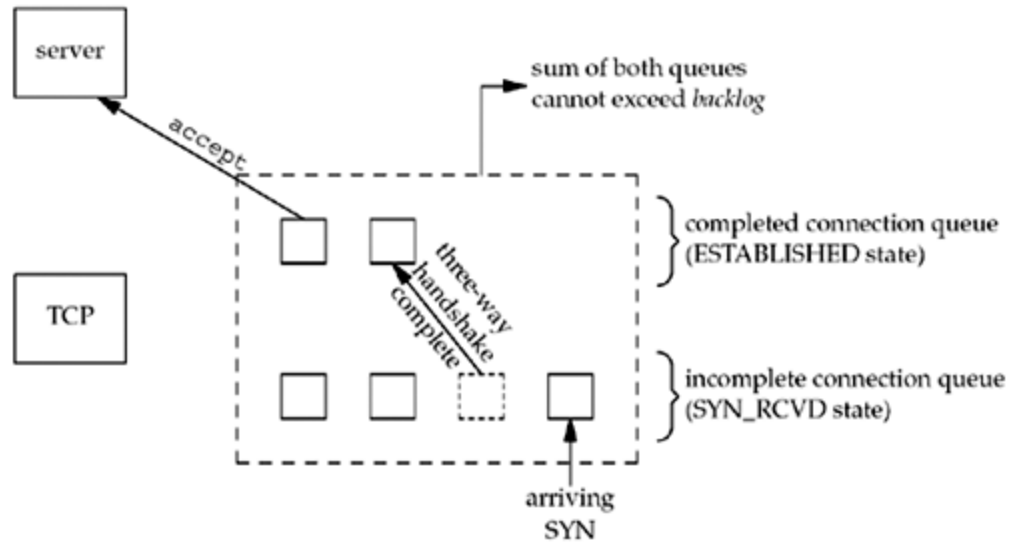



Figure 4.8. TCP three-way handshake and the two queues for a listening socket.



Accept

```
int newfd = accept(sockfd, (struct sockaddr *) & their_addr,  
                      &sin_size);
```

- Receive a new connection
 - ❖ their_addr contains the address of the remote host
- newfd is the socket descriptor for communications
 - ❖ Used like a regular socket, can read write close

Outline of a typical concurrent server

```
pid_t pid;
int  listenfd, connfd;

listenfd = Socket( ... );


/* fill in sockaddr_in{} with server's well-known port */
Bind(listenfd, ... );
Listen(listenfd, LISTENQ);

for ( ; ; ) {
    connfd = Accept (listenfd, ... );    /* probably blocks */

    if( (pid = Fork()) == 0) {
        Close(listenfd);    /* child closes listening socket */
        doit(connfd);       /* process the request */
        Close(connfd);      /* done with this client */
        exit(0);            /* child terminates */
    }

    Close(connfd);          /* parent closes connected socket */
}
```

Summary of TCP client/server with two clients.



Summary of UDP client/server with two clients.




Figure 4.14. Status of client/server before call to `accept` returns.



Figure 4.15. Status of client/server after return from `accept`.




Figure 4.16. Status of client/server after `fork` returns.





Figure 4.17. Status of client/server after parent and child close appropriate sockets.



Server

```
int main(int argc, char **argv) {
    int sockfd = socket(PF_INET, SOCK_STREAM, 0);
    struct sockaddr_in my_addr;
    my_addr.sin_addr = INADDR_ANY;
    my_addr.sin_port = htons(atoi(argv[1]));
    my_addr.sin_family = AF_INET;
    bind(sockfd, (struct sockaddr *) my_addr, sizeof my_addr);
    listen(sockfd, 10);
    while (1) {
        struct sockaddr_in their_addr;
        int sin_size;
        int newfd = accept(sockfd, (struct sockaddr *) & their_addr,
            &sin_size);
        write(newfd, "Hello world!\n", 14);
        close(newfd);
    } }
}
```

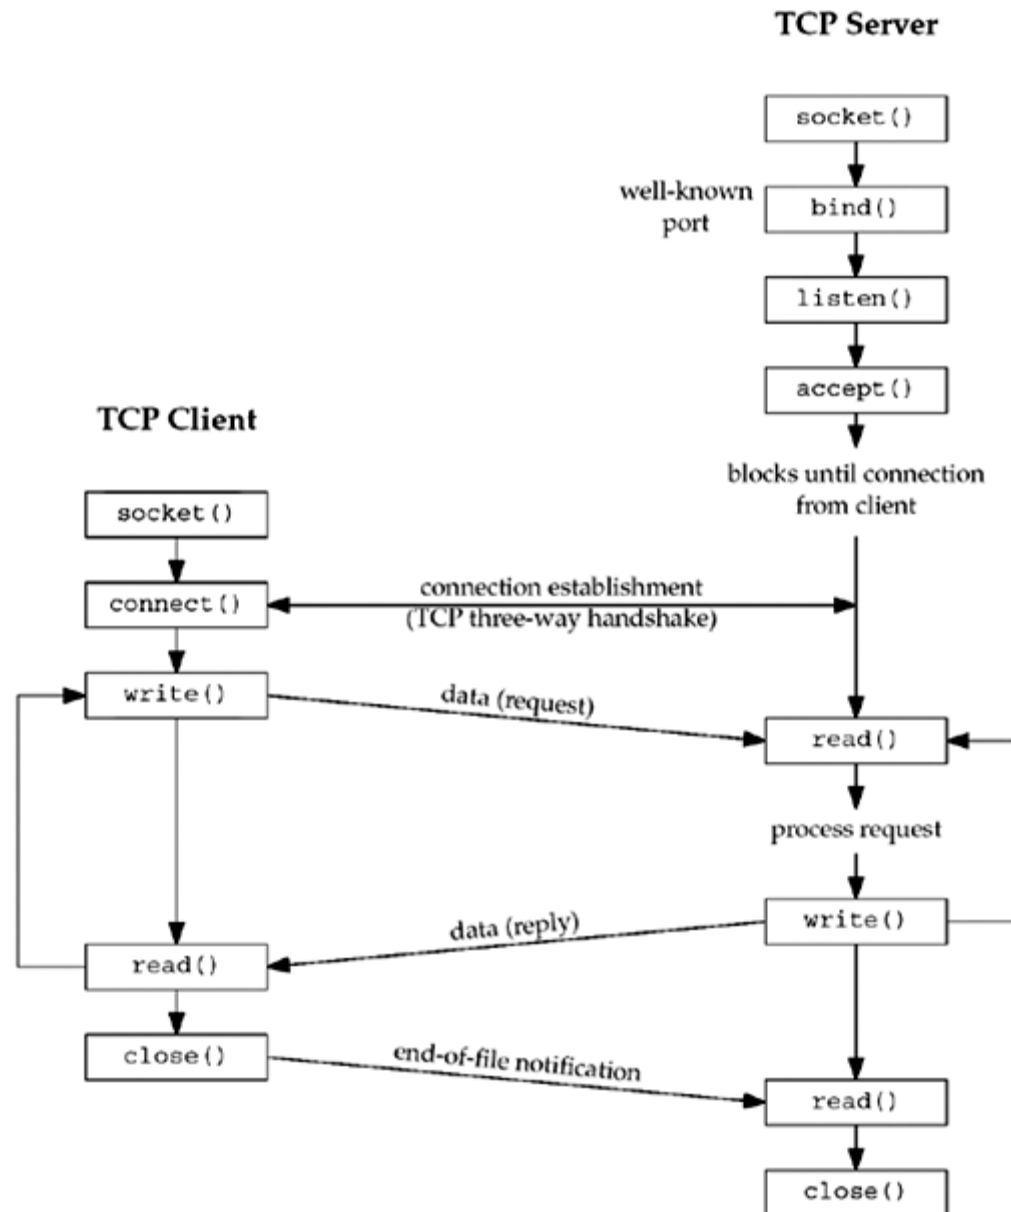



Figure 8.1. Socket functions for UDP client/server.

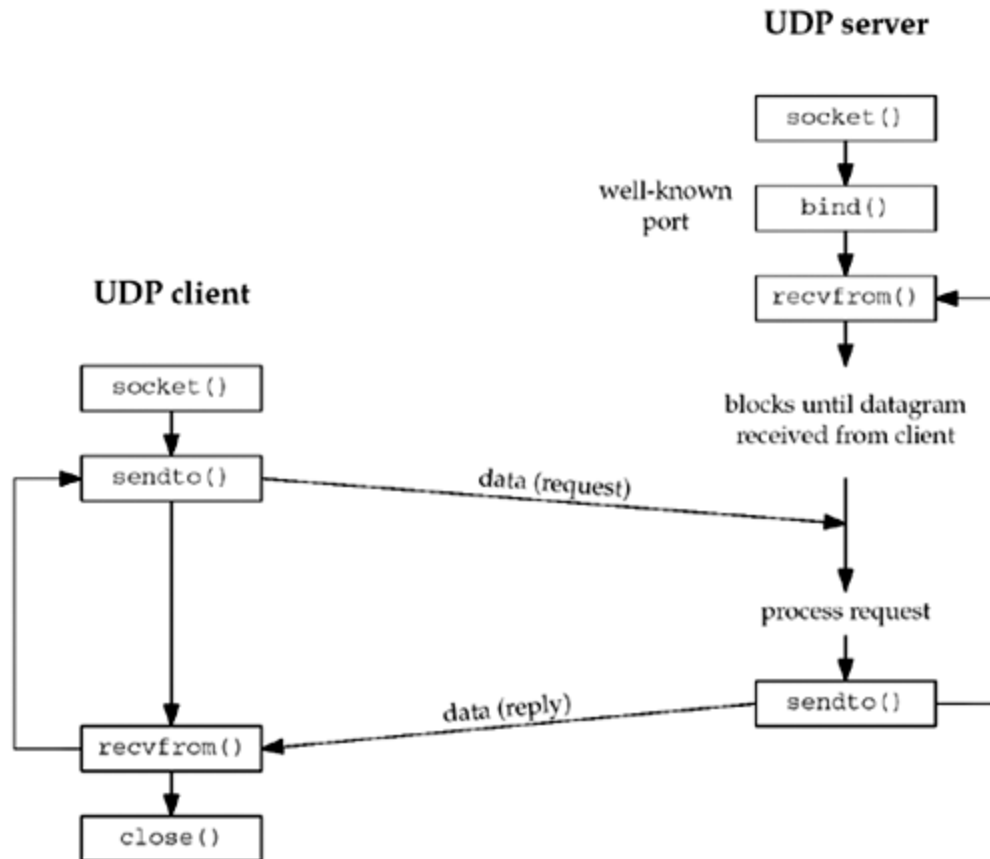



Figure 8.2. Simple echo client/server using UDP.



Support Slides

Programming Interfaces for network programs


UNIX Process & Socket
programming

Process Control

- ❑ Outline
- ❑ process management
- ❑ Fork() and Exec() systems calls
- ❑ error handling: perror

Process Trees

- ❑ Only an existing process can create a new process
- ❑ Parent-Child relations



What is Process?

- ❑ An entry in the kernel's process table
- ❑ Most common unit of execution
- ❑ Execution state
- ❑ Machine instructions, data and environment

Unix Process Model

- Simple and powerful primitives for process creation and initialization.
 - ❖ **fork** syscall creates a *child* process as (initially) a clone of the parent
 - ❖ parent program runs in child process to set it up for **exec**
 - ❖ child can **exit**, parent can **wait** for child to do so.
- Rich facilities for controlling processes by asynchronous *signals*.
 - ❖ notification of internal and/or external events to processes or groups
 - ❖ the look, feel, and power of interrupts and exceptions
 - ❖ default actions: stop process, kill process, dump core, no effect
 - ❖ user-level handlers

Process State

□ Process States

- ❖ **Running:** *the process is currently using CPU*
- ❖ **Runnable:** *the process can make use of the CPU*
- ❖ **Sleeping:** *the process is waiting for an event, e.g. I/O, to occur*
- ❖ **Suspended:** *the process is frozen by a signal; another signal is needed to wake up*
- ❖ **Idle:** *the process is being created by a fork() and not yet runnable*
- ❖ **Zombified:** *the process ended but not returned its exit code to its parent; a process remain zombie until its parent accepts its return code via the wait() system call*

Unix Process Creation

- ❑ Every process, except process 0, is created by the `fork()` system call
 - ❖ `fork()` allocates entry in process table and assigns a unique PID to the child process
 - ❖ child gets a copy of process image of parent: both child and parent are executing the same code following `fork()`
 - ❖ but `fork()` returns the PID of the child to the parent process and returns 0 to the child process

Unix Process Control

```
int pid;  
int status = 0;  
  
if (pid = fork()) {  
    /* parent */  
    .....  
    pid = wait(&status);  
} else {  
    /* child */  
    .....  
    exit(status);  
}
```

*The **fork** syscall returns a zero to the child and the child process ID to the parent.*

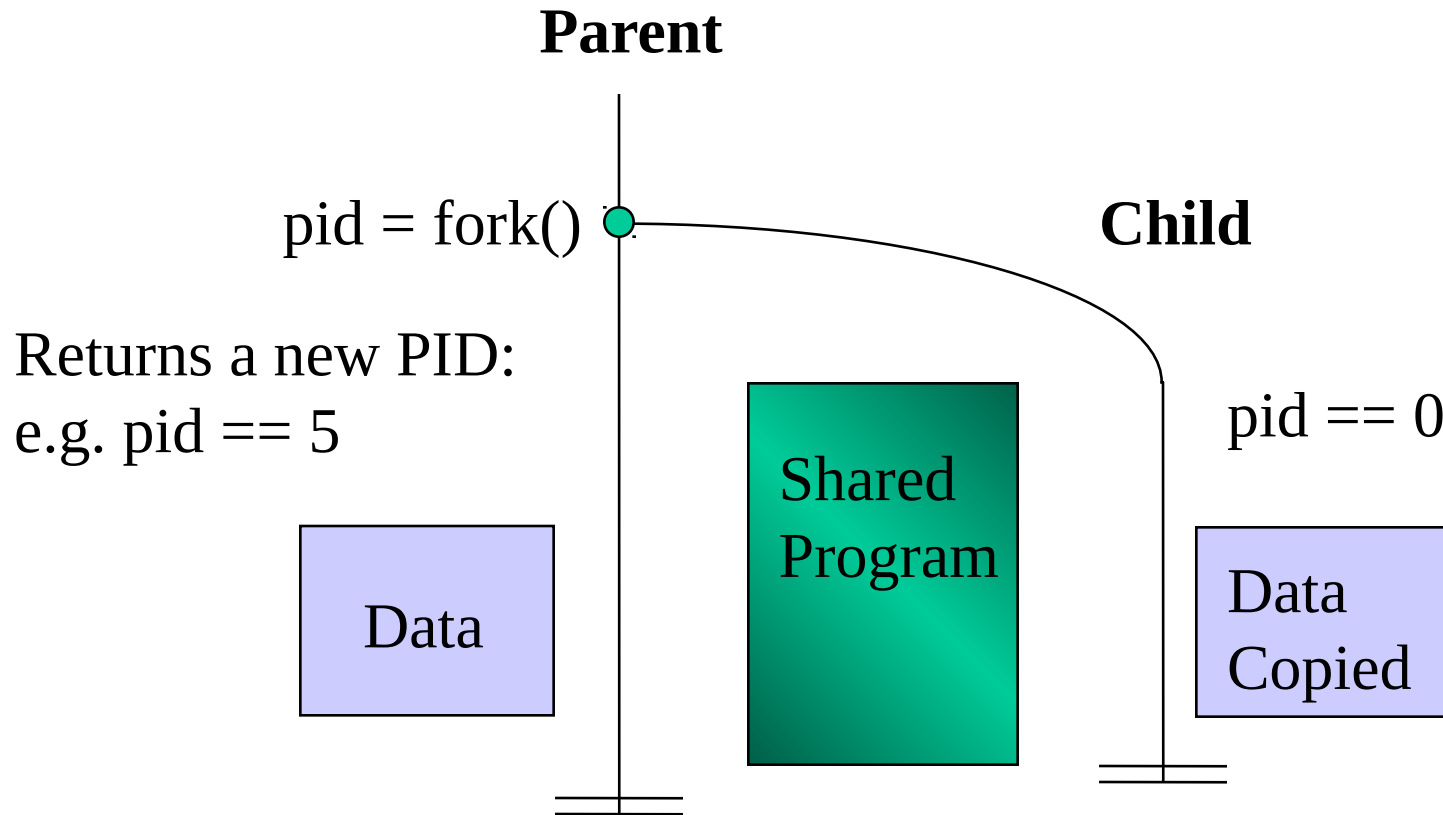
***Fork** creates an exact copy of the parent process.*

*Parent uses **wait** to sleep until the child exits; wait returns child pid and status.*

***Wait** variants allow wait on a specific child, or notification of stops and other signals.*

*Child process passes status back to parent on **exit**, to report success/failure.*

fork() as a diagram



Child Discipline

- ❑ After a *fork*, the parent program has complete control over the behavior of its child process.
- ❑ The child inherits its execution environment from the parent...but the parent *program* can change it.
 - ❖ sets bindings of file descriptors with *open*, *close*, *dup*
 - ❖ *pipe* sets up data channels between processes
- ❑ Parent program may cause the child to execute a different program, by calling *exec** in the child context.

Fork usage Key Points

- Parent and child both run same code
- Distinguish parent from child by return value from fork
- Start with same state, but each has private copy
- Including shared output file descriptor
- Relative ordering of their print statements undefined

Example using fork()

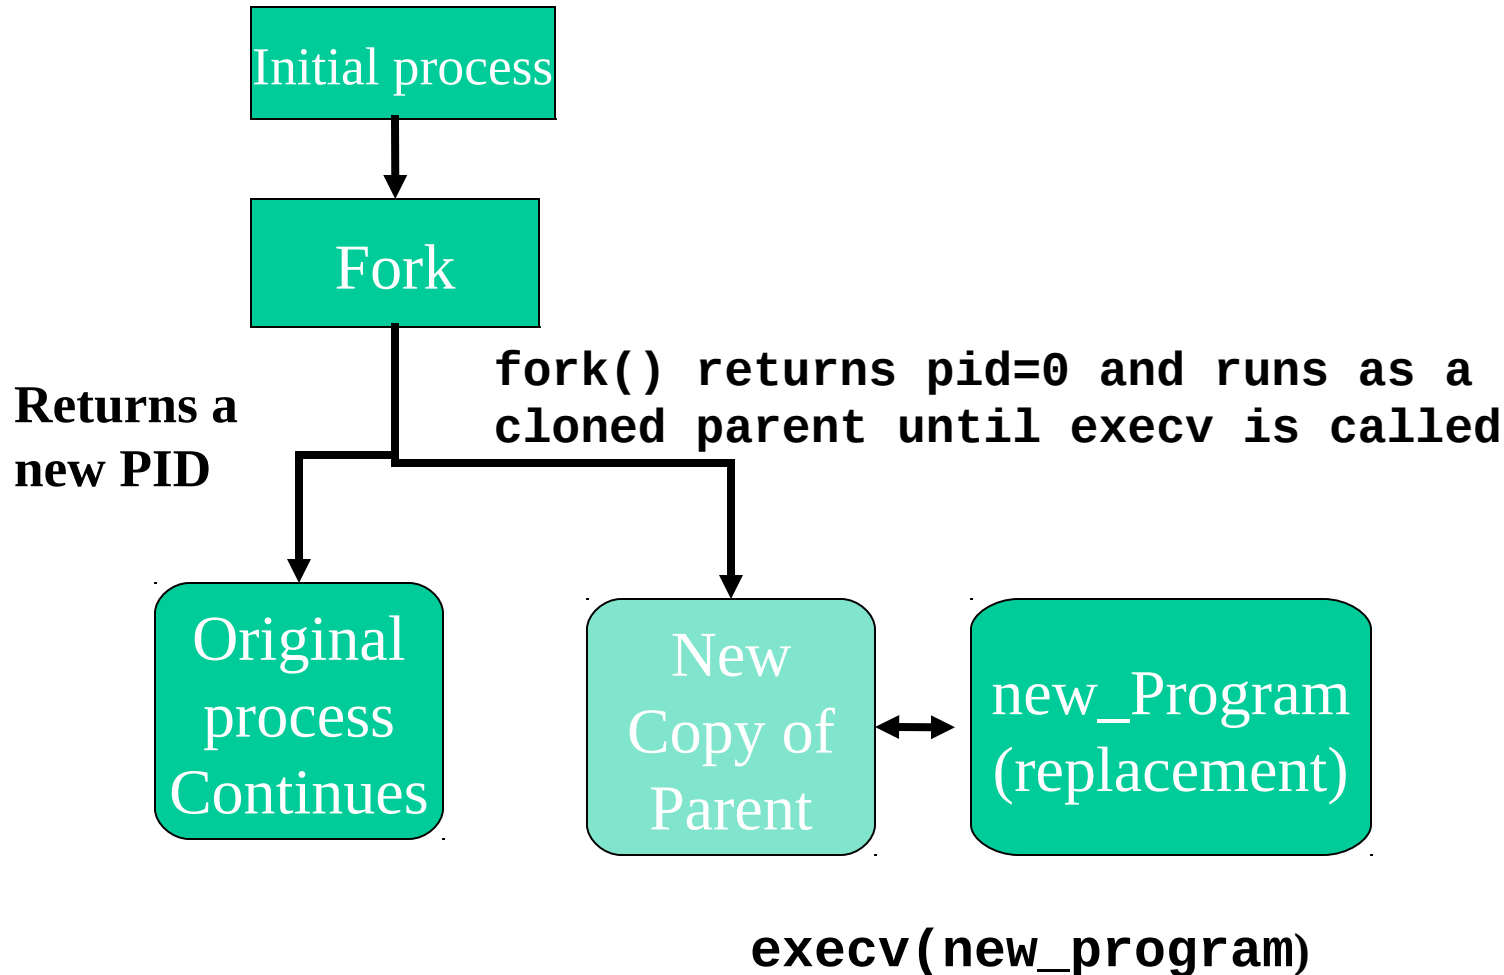
```
void fork1()  
{  
    int x = 1;  
    pid_t pid = fork();  
    if (pid == 0)  
    {  
        printf("Child has x = %d\n", ++x);  
    }  
    else  
    {  
        printf("Parent has x = %d\n", --x);  
    }  
    printf("Bye from process %d with x = %d\n", getpid(), x);  
}
```


Exec, Execve, etc.

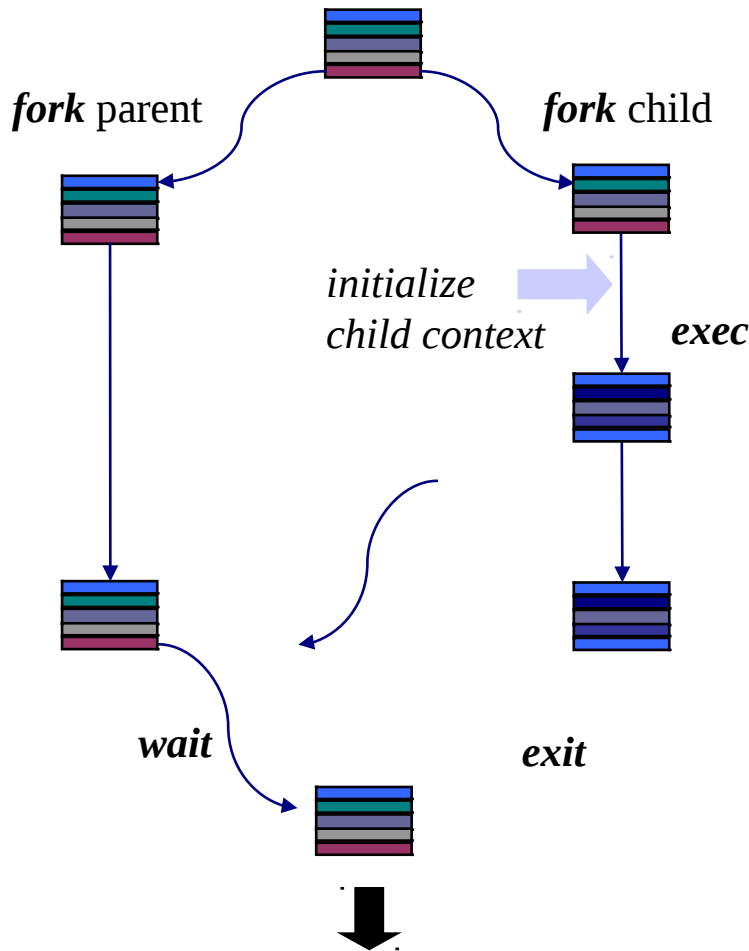
- ❑ Children should have lives of their own.
- ❑ Exec* “boots” the child with a different executable image.
 - ❖ parent program makes exec* syscall (in forked child context) to run a program in a new child process
 - ❖ exec* overlays child process with a new executable image
 - ❖ restarts in user mode at predetermined entry point (e.g., *crt0*)
 - ❖ no return to parent program (it’s gone)
 - ❖ arguments and environment variables passed in memory
 - ❖ file descriptors etc. are unchanged

fork() and execv()

❑ `execv(new_program, argv[])`



Fork/Exec/Exit/Wait Example



```
int pid = fork();
```

Create a new process that is a clone of its parent.

```
exec*("program" [, argvp, envp]);
```

Overlay the calling process virtual memory with a new program, and transfer control to it.

```
exit(status);
```

Exit with status, destroying the process.

```
int pid = wait*(&status);
```

Wait for exit (or other status change) of a child.


exec(..) Family

- There are 6 versions of the exec function, and they all do about the same thing: they replace the current program with the text of the new program. Main difference is how parameters are passed.

Versions of Exec

- ❑ `int execl(const char *path, const char *arg, ...);`
- ❑ `int execlp(const char *file, const char *arg, ...);`
- ❑ `int execl(const char *path, const char *arg`
`, ..., char *const envp[]);`
- ❑ `int execv(const char *path, char *const argv[]);`
- ❑ `int execvp(const char *file, char *const argv[]);`
- ❑ `int execve(const char *filename, char *const argv`
`[], char *const envp[]);`

exec(..) Family



Execve Example

```
int pid;
int status = 0;

if (pid = fork()) {
    /* parent */
    .....
    pid = wait(&status);
} else {
    /* child */
    execve(const char *filename,          char *const argv [],
           char *const envp[]);
    // exec shouldn't return
    return(ERROR);
}
```

Exit Usage

- ❑ `void exit(intstatus)`
- ❑ exits a process
- ❑ Normally return with status 0
- ❑ `atexit()` registers functions to be executed upon exit

```
void cleanup(void)  
{  
    printf("cleaningup\n");  
}
```

```
void fork6()  
{  
    atexit(cleanup);  
    fork();  
    exit(0);  
}
```


Error Handling

- ❑ Should always check return code of system calls
 - ❖ Not only for 5 marks in your lab!!! ;-)
 - ❖ There are subtle ways that things can go wrong
 - ❖ Use the status info kernel provides us
- ❑ Approach in this class: Wrappers
- ❑ Different error handling styles
 - ❖ Unix-Style
 - ❖ Posix-Style

Perror usage

- ❑ Prints system Error Message
- ❑ Describes last error encountered during a call to system or library function.

```
char *bufptr;  
size_t szbuf;  
if ((bufptr = malloc(szbuf)) == NULL)  
{  
    perror("malloc");  
    exit(2);  
}
```